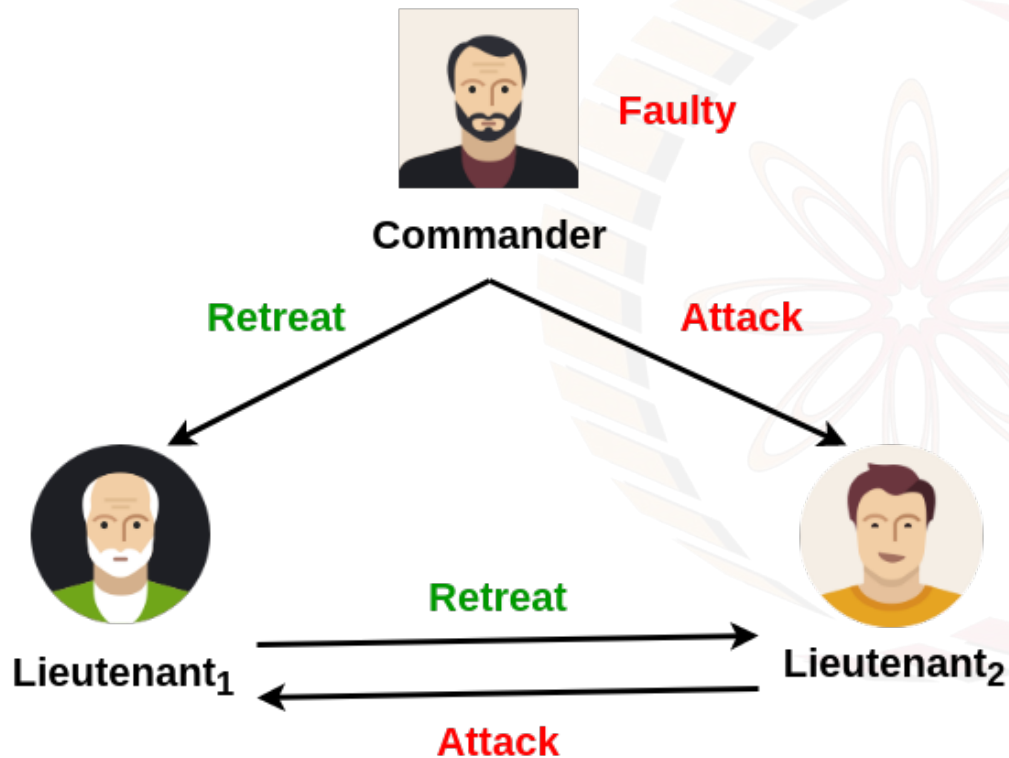



Three Byzantine Generals Problem: Commander Faulty



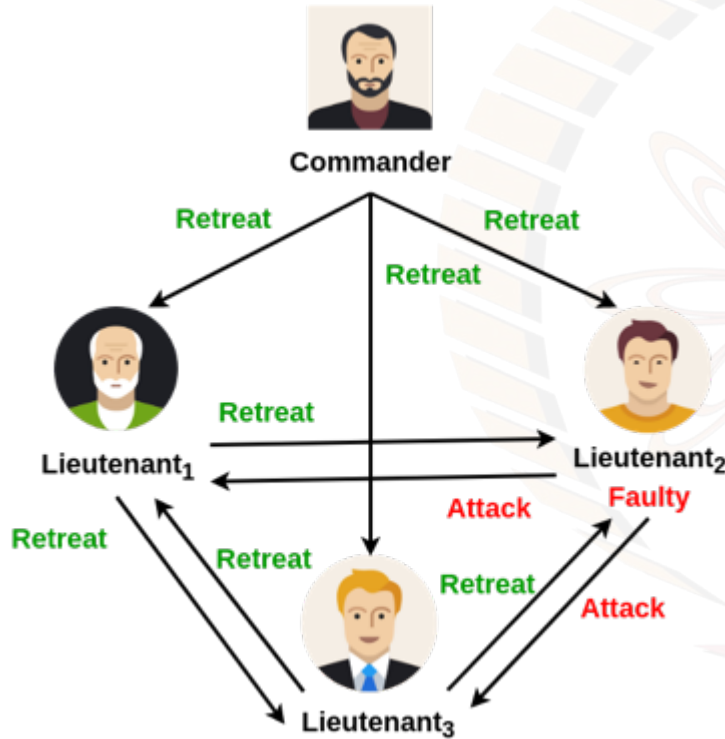
- Lieutenant₁ received **differing message**
- By integrity condition, both Lieutenants conclude with Commander's message
- This contradicts the agreement condition
- No solution possible for three generals including one faulty

Four Byzantine Generals Problem: Lieutenant Faulty




- Round 1:
 - Commander sends a message to each of the Lieutenants
- Round 2:
 - Lieutenant₁ and Lieutenant₃ correctly echo the message to others
 - Lieutenant₂ **incorrectly** echoes to others

Four Byzantine Generals Problem: Lieutenant Faulty



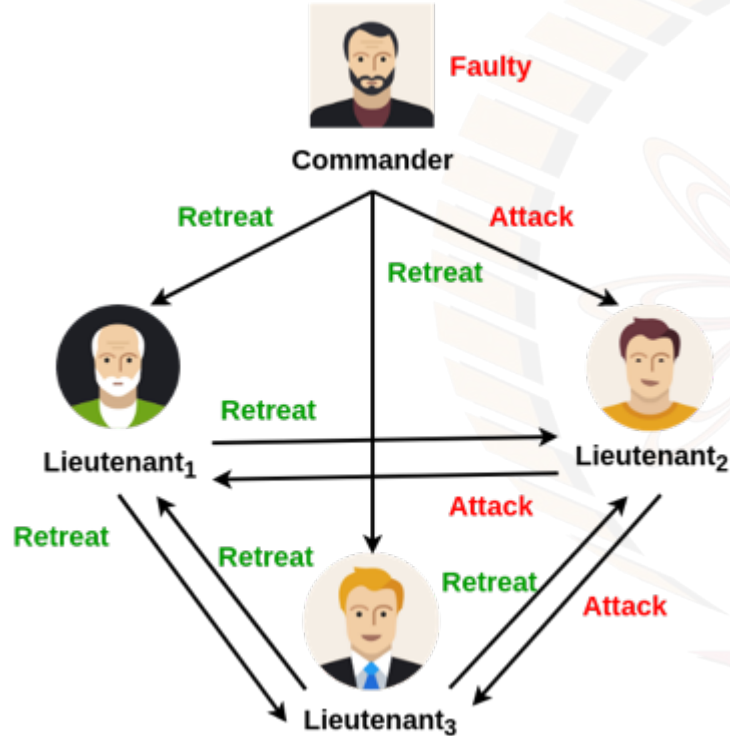
- Lieutenant₁ decides on $\text{majority}(\text{Retreat}, \text{Attack}, \text{Retreat}) = \text{Retreat}$
- Lieutenant₃ decides on $\text{majority}(\text{Retreat}, \text{Retreat}, \text{Attack}) = \text{Retreat}$

Four Byzantine Generals Problem: Commander Faulty




- Round 1:
 - Commander sends differing message to Lieutenants
- Round 2:
 - Lieutenant₁, Lieutenant₂ and Lieutenant₃ correctly echo the message to others

Four Byzantine Generals Problem: Commander Faulty




- Lieutenant₁ decides on $\text{majority}(\text{Retreat}, \text{Attack}, \text{Retreat}) = \text{Retreat}$
- Lieutenant₂ decides on $\text{majority}(\text{Attack}, \text{Retreat}, \text{Retreat}) = \text{Retreat}$
- Lieutenant₃ decides on $\text{majority}(\text{Retreat}, \text{Retreat}, \text{Attack}) = \text{Retreat}$

Byzantine Generals Model




- N number of process with at most f faulty
i.e. we have $2f + 1$ lieutenants and in case commandar is faulty concensus is reached for the value which commandar sent to the majority of the persons.
- Receiver always knows the identity of the sender
so closed
- Fully connected
- Reliable communication medium
- Synchronous system

Lamport-Shostak-Pease Algorithm



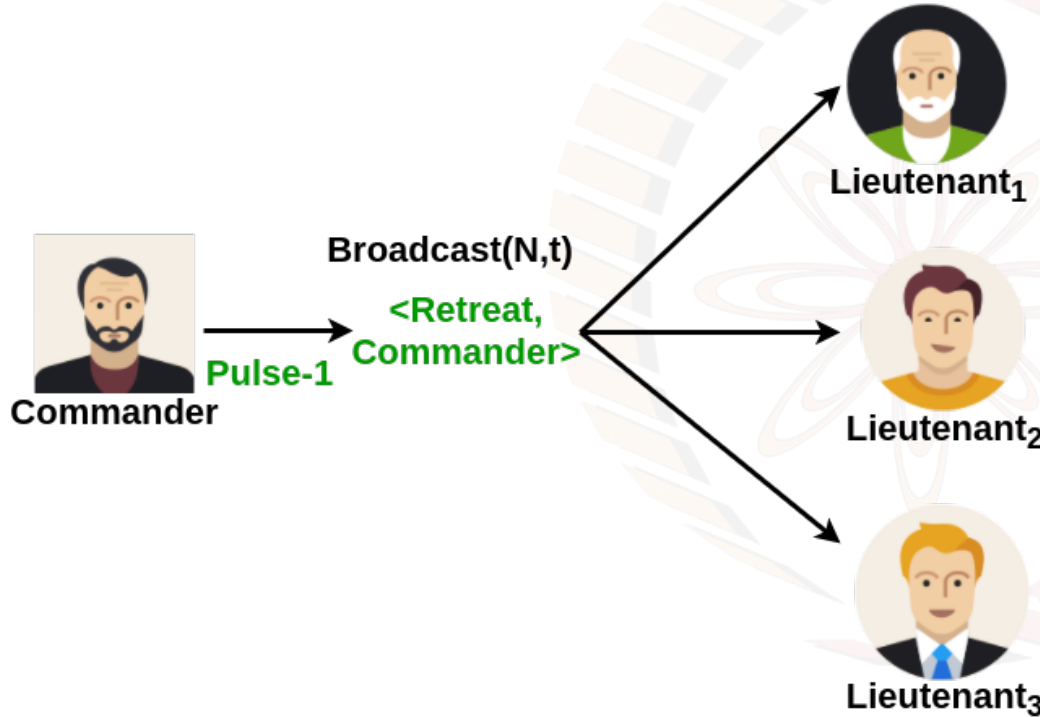
- **Base Condition:**
Broadcast($N, t=0$)
 - N : number of processes
 - t : algorithm parameter
- Commander decides on its own value

Lamport-Shostak-Pease Algorithm



- Base Condition:
Broadcast($N, t=0$)
 - N : number of processes
 - t : algorithm parameter
- Lieutenants decision by sender matching

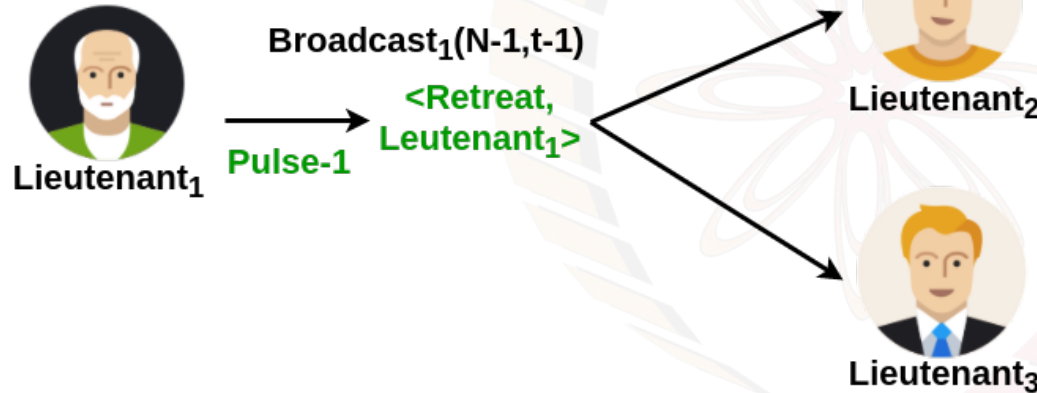
Lamport-Shostak-Pease Algorithm



- **General Condition:**
Broadcast(N,t)
 - N: number of processes
 - t: algorithm parameter
- Only commander sends to all lieutenants



Lamport-Shostak-Pease Algorithm



Now decide on majority voting

- **General Condition:**
Broadcast(N,t)
 - N: number of processes
 - t: algorithm parameter
- All lieutenants broadcast their values to the other lieutenants except the senders



Practical Byzantine Fault Tolerant

- Why Practical?
 - Ensures safety over an asynchronous network (not liveness!)
 - Byzantine Failure
 - Low overhead



Practical Byzantine Fault Tolerant Model

- A state machine is replicated across different nodes

This $3f + 1$ includes primary

- $3f + 1$ replicas are there where f is the number of faulty replicas
- The replicas move through a successions of configurations, known as *views*
- One replica in a *view* is *primary* and others are *backups*




Practical Byzantine Fault Tolerant Model

- Views are changed when a *primary* is detected as faulty
- Every view is identified by a unique integer number v
- Only the messages from the current views are accepted



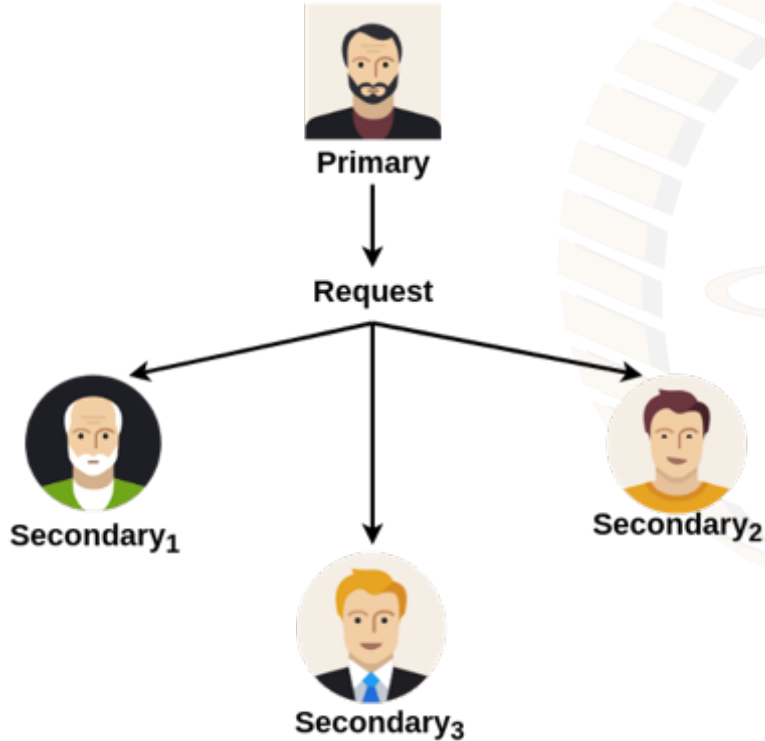
Practical Byzantine Fault Tolerant Algorithm



- A client sends a request to invoke a service operation to the primary

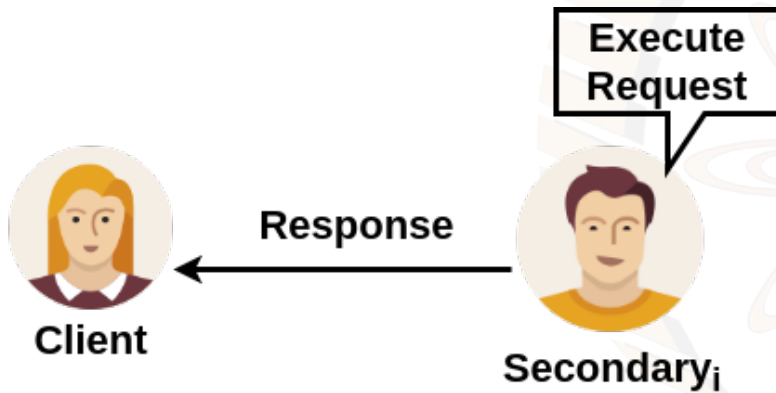


Practical Byzantine Fault Tolerant Algorithm



- The primary multicasts the request to the backups


Practical Byzantine Fault Tolerant Algorithm



- Backups execute the request and send a reply to the client



Practical Byzantine Fault Tolerant Algorithm




- The client waits for $f + 1$ replies from different backups with the same result
 - f is the maximum number of faulty replicas that can be tolerated

Three Phase Commit Protocol - Pre-Prepare

- **Pre-prepare:** Primary assigns a sequence number n to the request and multicast a message $\langle \langle PRE - PREPARE, v, n, d \rangle_{\sigma_p}, m \rangle$ to all the backups
 - v is the current view number
 - n is the message sequence number
 - d is the message digest
 - σ_p is the private key of primary - works as a digital signature
 - m is the message to transmit



Three Phase Protocol



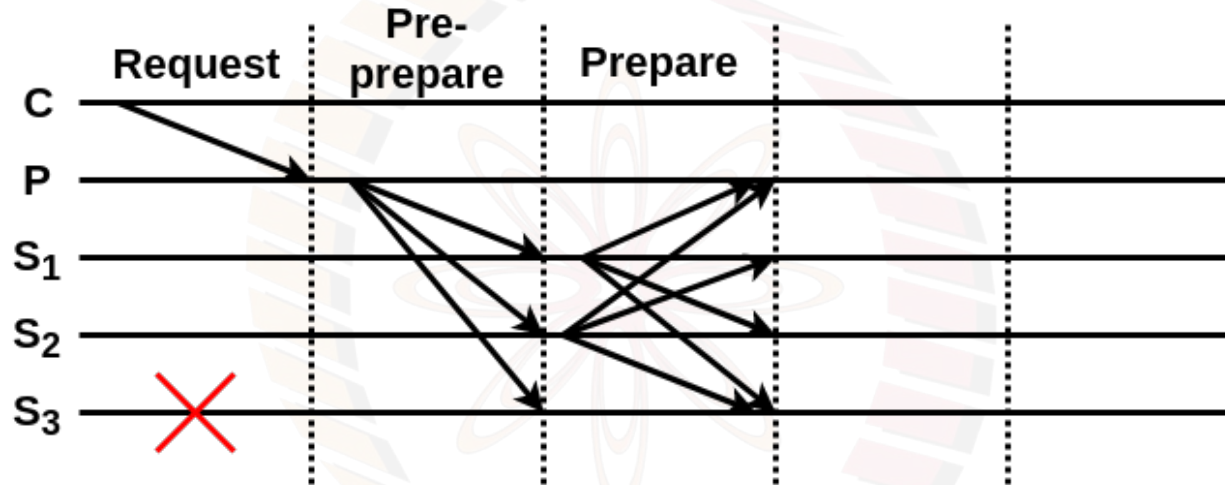
- Pre-prepare:
 - Acknowledge the request by a unique sequence number

Three Phase Commit Protocol - Pre-Prepare

- Pre-prepare messages are used as a proof that request was assigned sequence number n is the view v
- A backup accepts a pre-prepare message if
 - The signature is correct and d is the digest for m
 - The backup is in view v
 - It has not received a different PRE-PREPARE message with sequence n and view v with a different digest
 - The sequence number is within a threshold



Three Phase Protocol



- Prepare:
 - Replicas agree on the assigned sequence number



Three Phase Commit Protocol - Prepare

- If the backup accepts the PRE-PREPARE message, it enters prepare phase by multicasting a message $\langle PREPARE, v, n, d, i \rangle_{\sigma_i}$ to all other replicas
- A replica (both primary and backups) accepts prepare messages if
 - Signatures are correct
 - View number equals to the current view
 - Sequence number is within a threshold



Three Phase Commit Protocol

- Pre-prepare and prepare ensure that non-faulty replicas guarantee on a total order for the requests within a view
- Commit a message if
 - $2f$ prepares from different backups matches with the corresponding pre-prepare
 - You have total $2f + 1$ votes (one from primary that you already have!) from the non-faulty replicas



Three Phase Commit Protocol

- Why do you require $3f + 1$ replicas to ensure safety in an asynchronous system when there are f faulty nodes?
 - If you have $2f + 1$ replicas, you need all the votes to decide the majority - boils down to a synchronous system
 - You may not receive votes from certain replicas due to delay, in case of an asynchronous system
 - $f + 1$ votes do not ensure majority, may be you have received f votes from Byzantine nodes, and just one vote from a non-faulty node (note Byzantine nodes can vote for or against - You do not know that a priori!)

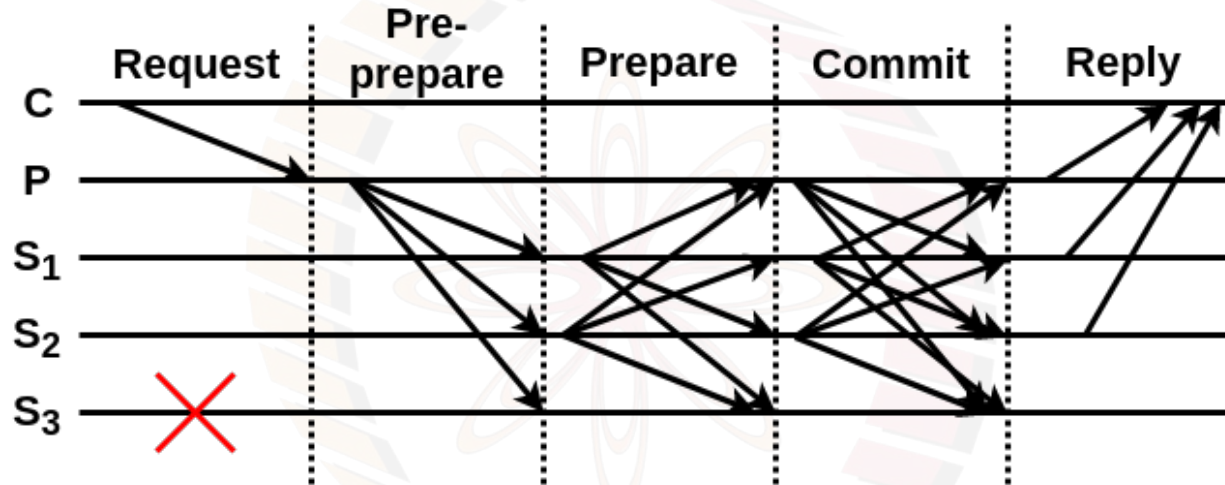


Three Phase Commit Protocol - Commit

- Multicast $\langle COMMIT, v, n, d, i \rangle_{\sigma_i}$ message to all the replicas including primary
- Commit a message when a replica
 - Has sent a commit message itself
 - Has received $2f + 1$ commits (including its own)



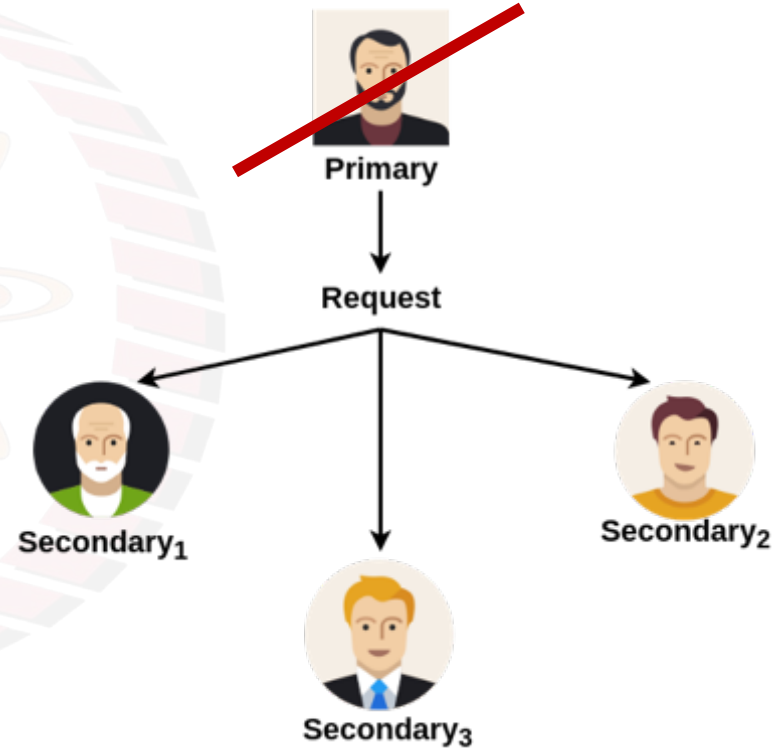
Three Phase Protocol



- Commit:
 - Establish consensus throughout the views

View Change

- What if the **primary** is **faulty**??
 - non-faulty replicas detect the fault
 - replicas together start view change operation



View Changes

- View-change protocol provides **liveness**
 - Allow the system to make progress when primary fails
- If the primary fails, backups will not receive any message (such as PRE_PREPARE or COMMIT) from the primary
- View changes are triggered by timeouts
 - Prevent backups from waiting indefinitely for requests to execute




View Change

- Backup starts a timer when it receives a request, and the timer is not already running
 - The timer is stopped when the request is executed
 - Restarts when some new request comes
- If the timer expires at view v
 - Backup starts a view change to move the system to view $v + 1$



View Change

- On timer expiry, a backup stops accepting messages except
 - Checkpoint
 - View-change
 - New-View



View Change

- Multicasts a $\langle VIEW_CHANGE, v + 1, n, \mathcal{C}, \mathcal{P}, i \rangle_{\sigma_i}$ message to all replicas
 - n is the sequence number of the last stable checkpoint s known to i
 - \mathcal{C} is a set of $2f + 1$ valid checkpoint messages proving the correctness of s
 - \mathcal{P} is a set containing a set \mathcal{P}_m for each request m that prepared at i with a sequence number higher than n
 - Each set \mathcal{P}_m contains a valid pre-prepare message and $2f$ matching



View Change

- The new view is initiated after receiving $2f$ view change messages
- The view change operation takes care of
 - Synchronization of checkpoints across the replicas
 - All the replicas are ready to start at the new view $v + 1$



Correctness

- **Safety:** The algorithm provides safety if all non-faulty replicas agree on the sequence numbers of requests that commit locally



Correctness

Liveness: To provide liveness, replicas must move to a new view if they are unable to execute a request

- A replica waits for $2f + 1$ view change messages and then starts a timer to initiate a new view (*avoid starting a view change too soon*)
- If a replica receives a set of $f + 1$ valid view change messages for views greater than its current view, it sends view change message (*prevents starting the next view change too late*)
- Faulty replicas are unable to impede progress by forcing frequent view change

